

# Mini Clinical Note & Coding Assistant (Node.js full-stack)

## Project: Mini Clinical Note & Coding Assistant (Node.js full-stack)

### Problem statement

Build a small web app that turns a short primary-care consultation transcript into:

1. a structured clinical note (SOAP),
2. suggested problem list & ICD-10 codes,
3. a draft billing hint (e.g., likely E/M level or CPT short-list), and
4. an AI-generated compliance banner explaining limitations/safety ( “not a medical device,” “for clinician review only” ).

No telephony or live diarization is required—assume you’ re given a short transcript (or upload a text file). Keep it lightweight, but production-minded.

---

### Requirements (must-have)

#### A. Stack & architecture

- **Single repo, Node.js throughout.**
  - Backend: Node.js + Express (or Fastify).
  - Frontend: React (CRA, Vite, or Next.js) with TypeScript preferred.
  - Storage: in-memory or a lightweight DB (SQLite/LowDB) for session history.
- **AI integration:** Call *one* LLM (of your choice) via REST from the backend only. No keys in the browser.
- **Deterministic interfaces:** Define clean API routes, return JSON with explicit schemas (Zod/TypeBox ok).

#### B. Core features

1. Transcript input

- Paste text or upload a `.txt` file (<= 5 KB). -- Can use some consultation examples on the internet or youtube, etc.
  - Show a read-only view with basic formatting (speaker labels optional).
- ## 2. Note generation (Documentation)
- Produce a **SOAP** note with: Subjective, Objective, Assessment, Plan.
  - Include a **problem list** extracted from the transcript with brief rationales (1–2 lines each).
- ## 3. Coding suggestion (Billing & Coding)
- Suggest **up to 3** ICD-10 codes ranked by relevance (code + description).
  - Provide a **billing hint**: either a likely **E/M level** or **up to 3 CPT codes** with one-line justification each.
  - Add **model self-confidence**: low/med/high per suggestion.
- ## 4. Compliance & guardrails
- Prepend a **compliance banner**: “Draft only; clinician review required; not a medical device; may be inaccurate.”
  - Implement **two simple guardrails**:
    - If transcript contains emergencies ( “chest pain,” “suicidal,” etc.), show a **high-risk flag** and require a user checkbox ( “I understand this is not a triage tool” ) before revealing suggestions.
    - **No definitive medical claims**: convert absolutes ( “will cure” ) to cautious language ( “may help” ) before displaying.
- ## 5. Traceability
- Show the **exact prompts** sent to the model (redact secrets).
  - Keep a **decision log** in the UI with the key steps (e.g., “extracted symptoms → mapped to problem list → mapped to ICD-10” ).
- ## 6. Research notes
- Add a small “**Research.md**” in the repo listing the public sources you used to understand ICD-10/CPT/E/M basics (links + 1-line takeaways). Aim for reputable sources.

## C. UX expectations

- Clean, minimal, keyboard-friendly UI.
  - A single page is fine; use cards/tabs for: Transcript | Note | Codes | Billing | Trace.
  - Provide a **copy** button for each section and **export JSON** of the whole result.
-

## Evaluation rubric (100 pts)

- **Problem understanding & scope fit (10)** – Delivers exactly what’s asked; reasonable assumptions documented.
  - **Frontend craft (20)** – Component design, state management, accessibility, and polish (loading/empty/error states).
  - **Backend/API quality (15)** – Clear endpoints, validation, error handling, and separation of concerns.
  - **AI engineering (20)** – Prompt design (few-shot or structured), output schemas, guardrails, failure modes, and sensible retries/timeouts.
  - **Clinical structure & coding logic (15)** – SOAP clarity, problem extraction quality, reasonable ICD-10/CPT/E/M hints with justifications.
  - **Security & compliance thinking (10)** – Secrets handled server-side, basic PII hygiene, disclaimers, and “no medical device” guardrails.
  - **Research & internet use (5)** – Thoughtful, cited sources in Research.md; evidence of reading up (not guesswork).
  - **DX & delivery (5)** – Clear README, run scripts, and a short demo GIF or Loom link.
- 

## Stretch goals (pick 1–2 if time allows; not required)

- **Simple RAG:** Allow the user to toggle “ground coding suggestions on uploaded ICD-10 CSV” and show which entries influenced the answer.
  - **Model comparer:** Compare two models (e.g., API A vs. API B) and show diff of SOAP and codes.
  - **Basic PHI scrub:** Before sending to the model, replace names/phones with placeholders and show a mapping table in the Trace tab.
  - **Unit tests:** 3–5 focused tests, especially around parsing/model output shaping.
- 

## Deliverables

- **Repo with:**
  - `README.md` (setup, run, design choices, assumptions, time spent).
  - `Research.md` (links + 1-line insights).
  - `docs/` (optional diagrams: sequence/flow).
  - **Demo:** short GIF or Loom ( $\leq 2$  min) walking through transcript → outputs → trace.

- **Working app:** `npm run dev` (or `pnpm/yarn`) to start both server and UI.
  - **One sample transcript:** add a short synthetic 1–2 minute primary-care interaction (no real PHI).
- 

## Suggested approach (time guide, not mandatory)

1. **Scaffold & data path (1–2h):** Set up Node/Express, React, env vars, and a simple POST `/analyze`.
  2. **Prompting & shaping (1–2h):** Design a single “analysis” prompt that returns a typed JSON (SOAP, problems, codes, risks). Use a schema validator.
  3. **Guardrails (0.5–1h):** Keyword flagging for emergencies + claim softener.
  4. **UI polish (1–2h):** Tabs/cards, copy/export, loading/errors.
  5. **Docs & demo (0.5–1h):** README, Research.md, 90-second walkthrough.
- 

## Acceptance criteria (functional)

- Upload or paste transcript → click **Analyze** → see SOAP, problems, ICD-10, billing hint, banner, and trace.
  - Refresh the page and the last result still appears (simple persistence).
  - Prompts & model responses visible in Trace, with PII placeholders if you implemented PHI scrub.
  - No API keys in client code; errors are human-readable.
- 

## Implementation hints

- **Prompting:** Ask the model to output **strict JSON**; validate it. Provide few-shot examples inline to stabilize structure.
- **Coding hints:** You may use simple **string matching + heuristics** (e.g., map “Type 2 diabetes” → E11.9); document your mapping sources in Research.md.
- **Risk words list:** start tiny (e.g., “chest pain”, “shortness of breath”, “suicidal”, “anaphylaxis”).
- **Compliance banner:** hard-code plus model-generated expansion (e.g., “This draft may miss context...”).

- **Testing:** Add at least one unhappy path (malformed JSON from model → show friendly error and raw text in Trace).
- 

## What we're assessing (why this works)

- **Front-end strength** via a focused but polished UI.
  - **Full-stack fundamentals** (auth not required, but server boundaries, validation, and error handling are).
  - **Applied AI skill**—prompting, schema enforcement, guardrails, basic safety/compliance thinking relevant to healthcare.
  - **Research & judgement**—how you interpret ICD-10/CPT/E/M at a lightweight level, and how you cite sources succinctly.
  - **Delivery**—clear docs, runnable project, thoughtful trade-offs.
- 

## Submission

Send a Git repo link or demo the project face to face if proceed to the next round.